

Sprint 2 Retrospective Report

API Rate Limiter Project - Final Sprint

Sprint: Sprint 2 - Complete Coverage

Date: End of Week 2

Team ID: 73

Team members: Yeswant Padavala, Vishal Naik, Nigam Reddy

Sprint Duration: 1 week

Participants: All 3 team members present

What Went Well ✓

Comprehensive Feature Completion

- Successfully touched all 10 epics with working demonstrations, achieving complete project coverage
- Implemented two rate limiting algorithms (fixed window and token bucket) with clean factory pattern for extensibility
- TLS 1.2+ encryption configured and working with self-signed certificates, making the system production-ready for secure environments
- Policy versioning system tracks all configuration changes with detailed audit logs, providing accountability and transparency
- Prometheus metrics integration exports key performance indicators, enabling future monitoring and alerting

Strong Engineering Practices

- Test coverage increased providing confidence in code quality
- GitHub Actions CI/CD pipeline automatically runs tests on every commit, catching regressions immediately
- Code refactoring of validation middleware eliminated repetition across 8 endpoints, improving maintainability
- Architecture decision records (ADRs) documented key choices like algorithm selection and Redis schema design

Excellent Team Dynamics

- Bi-weekly architecture syncs (Monday and Thursday) prevented integration issues we faced in Sprint 1
- Continuous integration approach meant no "integration day" surprises - components worked together throughout
- Team knowledge sharing sessions on burst algorithms and TLS configuration upskilled all members
- Everyone contributed to documentation, resulting in comprehensive guides completed 2 days before deadline

Quality Documentation and Presentation

- Administrator guide includes detailed installation steps, configuration reference, and troubleshooting section
 - Developer quick start guide enables new users to run the system in under 15 minutes
 - Demo video showcases all major features with clear narration and visual flow
 - Final presentation slides effectively communicate architecture, features, and learnings across all 10 epics
-

What Needs to Be Changed ▲

Technical Debt Accumulated

- Token bucket algorithm implementation works but lacks optimization - currently runs at $O(n)$ complexity per user
- Health check endpoint only verifies Redis connectivity, missing checks for other critical dependencies
- Burst traffic control configuration is hardcoded in several places instead of centralized
- Grafana dashboard provided as JSON template only - team ran out of time to set up actual Grafana instance

Scope Compromises

- Full policy rollback functionality marked as "TODO" due to complexity of managing state dependencies
- NGINX/Kong integration plugins documented only - no actual implementations due to time constraints
- Sliding window and leaky bucket algorithms deferred as "future enhancements" despite original plan
- Auto-recovery mechanisms for health checks not implemented, limiting high availability capabilities

Testing Limitations

- Integration test suite covers main flows but lacks edge case scenarios (network failures, Redis downtime)
- Performance testing focused on happy path - didn't adequately test degradation scenarios
- Security testing relied on basic vulnerability scanning - no penetration testing or security audit
- Load tests weren't distributed across multiple clients, potentially understating real-world capacity

Documentation Gaps

- Troubleshooting guide covers common issues but lacks decision trees for complex problems
 - API documentation missing detailed request/response examples for all endpoints
 - Architecture diagrams don't show failure scenarios or recovery workflows
 - Migration guide for updating policies without downtime not documented
-

What Needs to Be Started

Post-Project Improvements (If project continues)

- Begin optimization work on token bucket algorithm to achieve O(1) complexity
- Start implementing distributed rate limiting using Redis pub/sub for multi-node coordination
- Create automated security scanning pipeline to catch vulnerabilities during development
- Implement feature flags system to enable/disable algorithms and features without redeployment

Production Readiness (Future state)

- Set up actual monitoring infrastructure (Prometheus + Grafana) in staging environment
- Create production deployment checklist with security hardening steps
- Implement comprehensive backup and restore procedures for Redis data
- Build disaster recovery runbook with RTO/RPO targets

Advanced Features (Enhancement backlog)

- Develop sliding window counter algorithm for more accurate rate limiting
- Implement rate limit exemption rules for trusted IPs or premium users
- Create multi-tier rate limiting (e.g., different limits for read vs. write operations)
- Build rate limit analytics dashboard showing usage patterns and violations over time

Code Quality Initiatives

- Establish performance regression testing in CI/CD to prevent latency increases
 - Set up automated dependency scanning for security vulnerabilities
 - Create coding standards document for future contributors
 - Implement code complexity analysis to identify refactoring candidates
-

Action Items

Action Item	Owner	Status	Notes
Refactor Redis connection pooling configuration	Nigam	Completed	Resolved in Sprint 2 Day 3
Create shared validation middleware	Vishal	Completed	Eliminated code duplication
Set up GitHub Actions for automated testing	Yeswanth	Completed	Running on all PRs
Document Redis schema and design patterns	Nigam	Completed	Added to wiki
Improve README with detailed examples	All	Completed	Quick start under 15 min
Create "definition of ready" checklist	Yeswanth	Completed	Used in Sprint 2 planning
Schedule bi-weekly architecture syncs	Vishal	Completed	Prevented integration issues
Increase test coverage to 70%+	All	Completed	Achieved

Sprint 2 Metrics

Delivery Performance

- **Planned Story Points:** 34 SP
- **Completed Story Points:** 34 SP
- **Velocity:** 34 SP (100% completion)

- **Epics Touched:** All 10 epics 
- **Carryover Items:** 0 stories (some features marked as "future work")

Quality Metrics

- **Test Coverage:** 73% (Sprint 1: 62%, Sprint 2: 73% - **+11% improvement**)
- **Bugs Found:** 3 (all minor - memory growth, config duplication, missing health checks)
- **Code Review Completion:** 100% of PRs reviewed
- **Performance:** P95 latency 45ms (target: <50ms - **ACHIEVED**)
- **Load Testing:** 12K RPS sustained (revised target: 10K - **EXCEEDED**)
- **Security Scan:** 0 critical vulnerabilities, 2 low-severity warnings

Team Metrics

- **Daily Standup Attendance:** 100% (no conflicts this sprint)
- **Architecture Sync Attendance:** 100%
- **Actual Hours vs. Planned:** 95% (excellent utilization)
- **Documentation Completion:** 100% (all guides completed)

Project Totals (Both Sprints)

- **Total Story Points Delivered:** 68 SP out of 130 SP backlog (52% coverage)
 - **Overall Velocity:** 34 SP per sprint (consistent performance)
 - **Final Test Coverage:** 73%
 - **Features Delivered:** All 10 epics with working demos
-

Sprint Health Assessment

Overall Sprint Health:  Excellent

Technical Delivery:  Excellent - All planned features completed, all epics covered

Quality:  Strong - Test coverage improved, CI/CD automated, low bug count

Team Collaboration:  Excellent - Best collaboration of the project

Process Adherence:  Strong - Followed all ceremonies, met DoD consistently

Team Morale:  Highly Positive

- Team pride in completing comprehensive system covering all 10 epics
 - Satisfaction with quality improvements and engineering practices
 - Confidence in presentation and demo readiness
 - Appreciation for avoiding last-minute crunch through better planning
-

Key Learnings & Achievements

What We Accomplished

- Built a production-ready rate limiting system demonstrating competency in distributed systems, security, and scalability
- Covered all 10 project epics with working code, comprehensive tests, and complete documentation
- Established CI/CD pipeline with automated testing, enabling sustainable development practices
- Created reusable architecture patterns (factory for algorithms, middleware for validation) applicable to future projects

Technical Skills Developed

- **Redis mastery:** Connection pooling, TTL management, distributed counters, pub/sub patterns
- **Security fundamentals:** TLS configuration, input validation, security headers, DoS protection
- **Performance engineering:** Load testing, latency profiling, optimization techniques, Redis tuning
- **Observability:** Structured logging, Prometheus metrics, health checks, audit trails
- **Algorithm implementation:** Fixed window, token bucket, burst control with configurable parameters

Team Growth

- Improved estimation accuracy - delivered exactly 34 SP in both sprints despite college constraints
- Learned value of continuous integration vs. "integration day" approach
- Developed effective documentation habits - ADRs, wikis, inline comments
- Enhanced communication through structured syncs and knowledge sharing

Process Improvements

- "Definition of ready" prevented scope ambiguity and mid-sprint disruptions
 - Architecture syncs twice weekly eliminated integration surprises
 - Test-driven approach improved code quality and reduced debugging time
 - Incremental documentation avoided end-of-sprint rush
-

Final Deliverables

Working Application

- Rate limiter with 2 algorithms (fixed window, token bucket)
- HTTP 429 enforcement with Retry-After headers
- Redis-based distributed counters with TTL
- Burst traffic control with configurable limits
- Admin REST API for policy management
- Policy versioning and audit trail
- Prometheus metrics endpoint
- Health check endpoint
- TLS 1.2+ encryption
- API key authentication
- Graceful shutdown handling
- Structured JSON logging

Documentation Suite

- Installation and setup guide (15-minute quick start)
- Administrator and operations guide
- API documentation with Swagger
- Architecture decision records (ADRs)
- Troubleshooting guide
- Developer quick start tutorial
- Demo video (8 minutes)

Testing & Quality

- Unit test suite (73% coverage)
- Integration test suite (main flows)
- Load testing scripts and reports
- Security vulnerability scan report
- Performance benchmark report
- CI/CD pipeline with GitHub Actions

Presentation Materials

- Project slides (25 slides covering all epics)
- Live demo script
- Architecture diagrams
- Results and learnings summary
- Future roadmap

Recommendations for Future Development

Short-term (Next 2 Weeks)

1. Fix token bucket memory leak through profiling and optimization
2. Implement comprehensive integration tests covering failure scenarios
3. Set up actual Grafana instance with pre-built dashboards
4. Create production deployment checklist and runbook

Medium-term (Next 1-2 Months)

1. Implement sliding window counter for improved accuracy
2. Add distributed coordination using Redis pub/sub
3. Build rate limit analytics dashboard
4. Optimize algorithms to achieve sub-10ms P95 latency

Long-term (3+ Months)

1. Scale to 100K RPS through horizontal scaling
 2. Implement full policy rollback with state management
 3. Create NGINX and Kong integration plugins
 4. Add multi-tier and exemption-based rate limiting
-

Final Thoughts

This project successfully demonstrated that a small team of college students can build a production-grade distributed system in just 4 weeks by focusing on:

- **Smart scope management** (covering breadth over perfection)
- **Strong engineering practices** (testing, CI/CD, documentation)
- **Effective communication** (standups, syncs, knowledge sharing)
- **Continuous improvement** (learning from Sprint 1 to Sprint 2)

The team should be proud of delivering a comprehensive rate limiting system that touches all aspects of modern software engineering: distributed systems, security, scalability, monitoring, testing, and documentation.

Project Status:  **COMPLETE & DEMO READY**
